

TIME SERIES ANALYSIS COURSEWORK 2020-2021

Juliette Limozin, CID: 01343907

Due 18/12/2020 at 4 pm

This is my own unaided work unless stated otherwise.

Packages used

```
1 import numpy as np
2 from scipy.fft import fft, fftshift
3 from scipy.linalg import toeplitz
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import random
7 random.seed(123)
8 pd.set_option('display.max_colwidth', 0)
```

Question 1

(a)

From lecture notes,

$$S(f) = \frac{\sigma_\epsilon^2}{|1 - \phi_{1,p}e^{-i2\pi f} - \dots \phi_{p,p}e^{-i2\pi fp}|^2}$$

```
1 def S_AR(f,phis,sigma2):
2     """
3     INPUT:
4         f: vector for frequencies to evaluate
5         phis: vector of [phi_1, .. phi_p] of AR(p) process
6         sigma2: variance of the white noise process
7     OUTPUT:
8         S: spectral density function of the AR(p) process evaluated at
9         frequencies in f
10    """
11    p = len(phis) #Determine p
12    A = 1 #initialise difference in the denominator
13    for i in range(1,p+1):
14        #Loop for each p to subtract from denominator
15        A = A - phis[i-1]*np.exp(-1j*2*np.pi*i*np.array(f))
16    S = sigma2/np.abs(A)**2 #compute spectral density function
17    return S
```

(b)

From results in problem sheets, we know that if the white noise process of a AR(2) process is Gaussian, then the AR(2) process is also Gaussian. Therefore we generate a Gaussian (normal) white noise process to simulate a Gaussian AR(2) process in the following code.

```
1 def AR2_sim(phis,sigma2,N):
2     """
3     INPUT:
4         phis: vector of [phi_1, phi_2] of Gaussian AR(2) process
5         sigma2: variance of the white noise process
6         N: desired length of output
7     OUTPUT:
8         X: vector of values of the generated AR(2) process, discarding
9         first 100 values
10    """
11    #generate Gaussian white noise process
```

```

12     et = np.random.normal(0, np.sqrt(sigma2), 100+N)
13     # initialise output X
14     X = np.zeros(100+N)
15     for t in range(2,100+N):
16         #loop to define each element X_t
17         X[t] = phis[0]*X[t-1]+phis[1]*X[t-2]+et[t]
18     X = X[100:] #discard first 100 values
19     return X

```

(c)

From lecture notes,

$$\hat{s}_\tau = \frac{1}{N} \sum_{t=1}^{N-|\tau|} X_t X_{t+|\tau|}$$

Assuming w.l.o.g that the τ s given are non-negative, we get the following code:

```

1 def acvs_hat(X,tau):
2     """
3     INPUT:
4         X: vector of time series
5         tau: vector of of values at which to estimate the autocovariance,
6             positive values
7     OUTPUT:
8         s: estimate of the autocovariance sequence at values in tau
9     """
10    #Define length of (X_1, ..., X_n)
11    N = len(X)
12    #initialise vector of autocovariance sequence
13    s = np.zeros(len(tau))
14    for i in range(len(tau)):
15        #loop to define each s_tau
16        s[i] = (1/N)*np.dot(X[:N-tau[i]], X[tau[i]:])
17    return s

```

Question 2

(a)

We get the following code using the formulas of the periodogram and direct spectral estimate with tapering from lectures notes.

```

1 def periodogram(X):
2     """
3     INPUT:
4         X: time series
5     OUTPUT:
6         S: periodogram of the time series at the Fourier frequencies
7     """
8     #Define length of the time series
9     N = len(X)
10    #Compute periodogram using fft
11    S = (1/N)*np.abs(fft(X))**2
12    return S
13 def direct(X):
14     """
15     INPUT:
16         X: time series
17     OUTPUT:
18         S: direct spectral estimate of the time series at the Fourier
19             frequencies using the Hanning taper
20     """
21    #Define length of the time series

```

```

22 N = len(X)
23 #Define the Hanning taper
24 t = np.array(range(1,N+1))
25 h = 0.5*(8/(3*(N+1)))**0.5 * (1-np.cos(2*np.pi*t/(N+1)))
26 #Compute the direct spectral estimate using the taper and fft
27 S = np.abs(fft(np.multiply(h,X)))**2
28 return S

```

(b)

We are given the roots of the AR(2) process at which we want to estimate the spectral density function. The roots are the solution to the equation $1 - \phi_{1,2}z - \phi_{2,2}z^2$. Then:

$$\begin{aligned}
 & \left(z - \frac{1}{r}e^{i2\pi f'}\right)\left(z - \frac{1}{r}e^{-i2\pi f'}\right) = 0 \\
 \implies & z^2 - \frac{z}{r}(e^{i2\pi f'} + e^{-i2\pi f'}) + \frac{1}{r^2} = 0 \\
 \implies & z^2 - \frac{z}{r}2\cos(2\pi/f') + \frac{1}{r^2} = 0 \\
 \implies & r^2z^2 - 2r\cos(2\pi/f')z + 1 = 0
 \end{aligned}$$

By matching coefficient terms, we get that

$$\begin{aligned}
 \phi_{1,2} &= 2r\cos(2\pi/f') = 2.0.95.\cos(\pi/4) = 0.95.\sqrt{2}; \\
 \phi_{2,2} &= -r^2 = -0.95^2
 \end{aligned}$$

Figure 1 shows the plots of the bias of the periodogram and direct spectral estimates at frequencies $[1/8, 2/8, 3/8]$ for different plots of sample sizes N , each simulated 10000 times. The x-axis is the log of N . To calculate the empirical bias at each frequency and estimate, I took the difference between the mean of the 10000 generated estimates and the true value of the spectral density at that frequency.

```

1 def question2b():
2     #vector of sample sizes to test
3     N = [16, 32, 64, 128, 256, 512, 1024, 2048, 4096]
4     #[phi_1,2, phi_2,2,]
5     phis = np.array([0.95*np.sqrt(2), -0.95**2])
6     #Frequencies at which to evaluate the bias of estimates
7     #true spectral density at these frequencies
8     f = [1/8, 2/8, 3/8]
9     S = S_AR(f, phis, 1)
10    #initialise matrices containing bias for each frequency and sample size
11    #and estimation method
12    bias_p = np.zeros((3,len(N)))
13    bias_d = np.zeros((3,len(N)))
14    for j in range(len(N)):
15        #loop for each sample size
16        S_p = np.zeros((3,10000))
17        S_d = np.zeros((3,10000))
18        for i in range(10000):
19            #generate time series
20            X_r = AR2_sim(phis,1,N[j])
21            #loop for 10000 realisations
22            Sp = periodogram(X_r)
23            Sd = direct(X_r)
24            S_p[:,i] = np.array([Sp[2**(j+1)], Sp[2**(j+2)], Sp[6*2**j]]).T
25            S_d[:,i] = np.array([Sd[2**(j+1)], Sd[2**(j+2)], Sd[6*2**j]]).T
26        #calculate bias
27        bias_p[:,j] = np.mean(S_p, axis = 1) - np.array(S).T
28        bias_d[:,j] = np.mean(S_d, axis = 1) - np.array(S).T
29    #plots
30    plt.figure(figsize = (15,20))
31    plt.subplot(311)
32    plt.plot(np.log(N), bias_p[0,:], label = 'Periodogram')
33    plt.plot(np.log(N), bias_d[0,:], label = 'Direct spectral estimate')

```

```

34 plt.legend()
35 plt.xlabel('Log N')
36 plt.ylabel('Bias')
37 plt.title('Bias of spectral estimators at frequency 1/8 for different' +
38           ' values of N')
39 plt.subplot(312)
40 plt.plot(np.log(N), bias_p[1,:], label = 'Periodogram')
41 plt.plot(np.log(N), bias_d[1,:], label = 'Direct spectral estimate')
42 plt.legend()
43 plt.xlabel('Log N')
44 plt.ylabel('Bias')
45 plt.title('Bias of spectral estimators at frequency 2/8 for different' +
46           ' values of N')
47 plt.subplot(313)
48 plt.plot(np.log(N), bias_p[2,:], label = 'Periodogram')
49 plt.plot(np.log(N), bias_d[2,:], label = 'Direct spectral estimate')
50 plt.legend()
51 plt.xlabel('Log N')
52 plt.ylabel('Bias')
53 plt.title('Bias of spectral estimators at frequency 3/8 for different' +
54           ' values of N')
55 plt.show()
56 return None

```

(c)

From lecture notes we know that because the periodogram and direct spectral estimates are consistent estimates of the spectral density function, as $N \rightarrow \infty$, $E(\hat{S}(f)) \rightarrow S(f)$. Therefore, $\lim_{N \rightarrow \infty} bias = \lim_{N \rightarrow \infty} E(\hat{S}(f)) - S(f) = 0$, which we can see in the plots of the bias of the estimates at frequencies $[1/8, 2/8, 3/8]$ in figure 1, as they all converge to 0.

Notice that for frequency $f = 1/8$, the biases are converging to 0 from below whereas for the other two frequencies they are converging from above. We can explain this by looking at the true spectral density function plotted in figure 2

As we can see, the spectral density moves towards $S(1/8)$ from below and towards $S(2/8)$ and $S(3/8)$ from above, which explains the equivalent behaviour for the bias.

We also note that direct spectral estimate using tapering has a lower bias than the periodogram. From lecture notes we know that this is because the tapering is specifically a bias reducing technique on the periodogram.

```

1 def question2c():
2     """
3     Code for question 2c
4     """
5     #Compute true spectral density
6     phis = np.array([0.95*2**0.5, -0.95**2])
7     f_l = [i/100 for i in range(100)]
8     S = S_AR(f_l, phis, 1)
9     #plots
10    plt.plot(f_l, S)
11    plt.vlines(1/8, 0, 160, color = 'r', linestyle = 'dashed', label = '1/8')
12    plt.vlines(2/8, 0, 160, color = 'g', linestyle = 'dashed', label = '2/8')
13    plt.vlines(3/8, 0, 160, color = 'b', linestyle = 'dashed', label = '3/8')
14    plt.legend()
15    plt.xlabel('frequency f')
16    plt.ylabel('S(f)')
17    plt.title('Plot of the true spectral density function')
18    plt.show()
19    return None

```

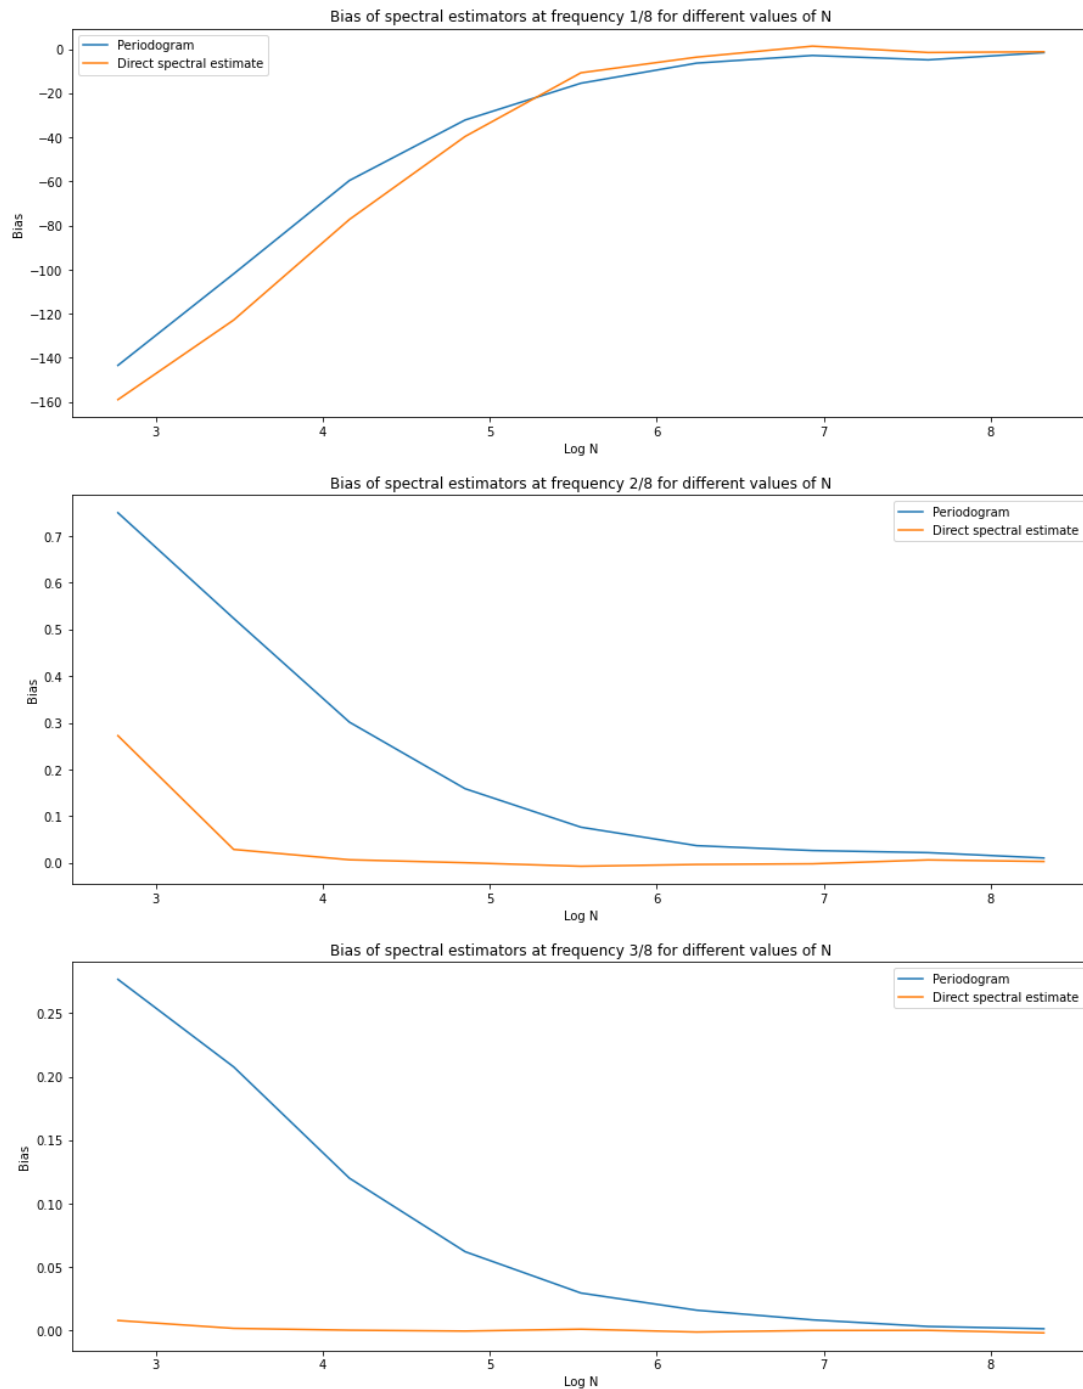


Figure 1: Plot of biases for question 2(b)

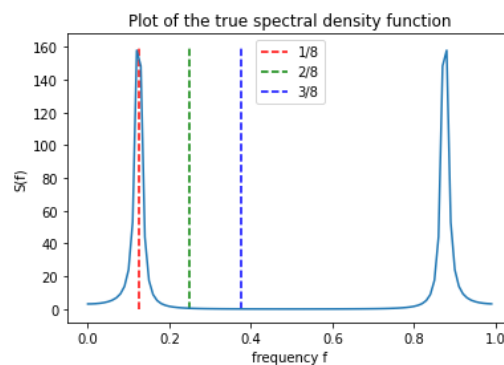


Figure 2: Plot of the true spectral density for question 2(c)

Question 3

(a)

Figure 3 shows the plots of the periodogram and the direct spectral estimate using the Hanning taper of my time series. The code is the same as `periodogram` and `direct` from question 2(a), but with an added `fftshift` to get the spectral density estimate on frequencies $[-1/2, 1/2]$.

```
1 def question3a(X):
2     """
3     Code for question 3a
4     INPUT:
5         X: my time series
6     """
7     N = len(X)
8     #compute periodogram with fftshift
9     S_p = (1/N)*np.abs(fftshift(fft(X)))**2
10    #Compute tapered direct spectral estimate with fftshift
11    t = np.array(range(1,N+1))
12    h = 0.5*(8/(3*(N+1)))**0.5 * (1-np.cos(2*np.pi*t/(N+1)))
13    S_d = np.abs(fftshift(fft(np.multiply(h,X))))**2
14    #Plots
15    f = [i/128 for i in range(-64,64)]
16    plt.figure(figsize = (20,10))
17    plt.subplot(121)
18    plt.plot(f, S_p)
19    plt.xlabel('frequencies f')
20    plt.ylabel('Periodogram')
21    plt.title('Periodogram of my time series at frequencies [-0.5,0.5]')
22    plt.subplot(122)
23    plt.plot(f, S_d)
24    plt.xlabel('frequencies f')
25    plt.ylabel('Direct spectral estimate')
26    plt.title('Direct spectral estimate of my time series using the Hanning taper'
27              + ' at frequencies [-0.5,0.5]')
28    plt.show()
29    return None
```




Figure 3: Plots of spectral density estimates for question 3(a)

(b)

The following is my code for fitting a AR(p) model using the Yule-Walker, Least Squares and approximate maximum likelihood methods, from formulas in lecture notes.

For the Least Squares fitting, we assume w.l.o.g that $F^T F$ is invertible.

```
1 def yule_walker(X,p):
2     """
3     INPUT:
4         X: time series of length 128
5         p: order of AR(p) process that we want to fit to X
6     OUTPUT:
7         phi_hat = (big gamma)^-1 * small gamma, yule-walker estimator of AR(p)
8         parameters phis
9         sigma2_hat: yule-walker estimator of variance of white noise process
10    """
11    s_hat = acvs_hat(X, list(range(0,p+1))) #estimate autocovariance sequence
12    gamma = s_hat[1:] #define small gamma = [shat_1, ..., shat_p]
13    #define big gamma, which is a symmetric Toeplitz matrix
14    Gamma = toeplitz(s_hat[:-1], s_hat[:-1])
15    #compute phi_hat and sigma2_hat
16    phi_hat = np.linalg.inv(Gamma).dot(gamma)
17    sigma2_hat = s_hat[0] - np.dot(phi_hat,s_hat[1:])
18    return phi_hat, sigma2_hat
19 def least_squares(X,p):
20     """
21     INPUT:
22         X: time series of length 128
23         p: order of AR(p) process that we want to fit to X
24     OUTPUT:
25         phis_hat = (F.T F)^-1 F.T X, least squares estimator of AR(p) parameters ↔
26         phis
27         sigma2_hat: lest squares estimator of variance of white noise process
28     """
29    N = len(X)
30    #initialise matrix F
31    F = np.zeros((N-p,p))
32    for i in range(p):
33        # Loop to define each column of F
34        F[:,i] = X[(p-i-1):(N-1-i)]
35    #Compute phi_hat
36    phi_hat = np.linalg.inv(F.T.dot(F)).dot(F.T).dot(X[p:])
37    #Compute sigma2_hat
38    sigma2_hat = (1/(N-2*p))*(X[p:]-F.dot(phi_hat)).T.dot(X[p:]-F.dot(phi_hat))
39    return phi_hat, sigma2_hat
40 def approx_mle(X,p):
41     """
42     INPUT:
43         X: time series of length 128
44         p: order of AR(p) process that we want to fit to X
45     OUTPUT:
46         phis_hat: least squares estimator of AR(p) parameters phis (which is
47         equal to Approximate MLE)
48         sigma2_hat: approximate maximum likelihood estimator of sigma2, which
49         is the lest squares estimator of sigma2 *(N-2p)/(N-p)
50     """
51    N = len(X)
52    #Get least squares estimators of phi and sigma2
53    phi_hat, sigma2_hat = least_squares(X, p)
54    #Transform least squares estimator of sigma2 into approximate MLE
55    sigma2_hat = sigma2_hat *(N-2*p)/(N-p)
56    return phi_hat, sigma2_hat
```

(c)

The AIC for each fitting method and process order is summarised in the table below.

	Yule-Walker	Least Squares	Approximate MLE
p = 1	83.457568	79.308525	78.296662
p = 2	67.210892	62.744916	60.696873
p = 3	42.897565	35.014882	31.905417
p = 4	10.712505	-22.692807	-26.889905
p = 5	11.667585	-23.526619	-28.838585
p = 6	12.558746	-19.564828	-26.019978
p = 7	13.945686	-16.908833	-24.536622
p = 8	15.673155	-13.065779	-21.896866
p = 9	16.615987	-11.939696	-22.006017
p = 10	18.121540	-11.858851	-23.193686
p = 11	19.796076	-7.995110	-20.633170
p = 12	21.759412	-7.187960	-21.165469
p = 13	22.142904	-5.101213	-20.456005
p = 14	23.842761	-3.255814	-20.027431
p = 15	25.429476	0.521702	-17.708102
p = 16	26.306425	2.717655	-17.013632
p = 17	26.731132	3.343002	-17.935132
p = 18	28.727868	5.515377	-17.357172
p = 19	30.279313	8.161516	-16.355375
p = 20	31.917098	12.468884	-13.744801

```

1 def question3c(X):
2     """
3     Code for question 3c
4     INPUT:
5         X: my time series
6     OUTPUT:
7         data_AIC: dataframe of AIC scores for each method and process order
8     """
9     #Initialise matrix of AICS
10    AIC = np.zeros((20,3))
11    N = len(X)
12    for p in range(1,21):
13        #Loop for computing AIC at each order p
14        _, s_y = yule_walker(X, p)
15        _, s_l = least_squares(X, p)
16        _, s_m = approx_mle(X, p)
17        sigma2s = np.array([s_y, s_l, s_m])
18        ps = np.array([p, p, p])
19        AIC[p-1,:] = 2*ps.T + N*np.log(sigma2s.T)
20    #Summarise in a pandas DataFrame
21    data_AIC = pd.DataFrame(data = AIC,
22                           index = ["p = " + str(i) for i in range(1,21)],
23                           columns = ['Yule-Walker',
24                                     'Least Squares', 'Approximate MLE'])
25    return data_AIC

```

(d)

We choose the order of model with the lowest AIC score for each fitting method.

Hence we choose $p = 4$ for the Yule-Walker model, and $p = 5$ for both the Least Squares and approximate maximum likelihood models.

The corresponding parameter estimates are shown in the table below, up to 4 decimal points.

	Yule-Walker	Least Squares	Approximate MLE
$\hat{\phi}$	[-1.5452, -1.2955, -1.0781, -0.4841]	[-1.7241, -1.6180, -1.4260, -0.7367, -0.0647]	[-1.7241, -1.6180, -1.4260, -0.7367, -0.0647]
$\hat{\sigma}^2$	1.0214	0.7696	0.7383

```

1 def question3d(X):
2     """

```



```

3 Code for question 3d
4 INPUT:
5     X: my time series
6 OUTPUT:
7     data_params: dataframe of p+1 parameters for each method
8     """
9 #Compute fitted parameters for chosen p
10 py, sy = yule_walker(X,4)
11 pl,sl = least_squares(X, 5)
12 pm,sm = approx_mle(X, 5)
13
14 #Summarise parameters into a pandas DataFrame
15 data_params = pd.DataFrame(data = [[ py, pl, pm],
16                                   [sy, sl, sm]],
17                             index = ['phi_hat', 'sigma2_hat'],
18                             columns = ['Yule-Walker',
19                                       'Least Squares', 'Approximate MLE'])
20
21 return data_params

```

(e)

Figure 4 shows the associated spectral density function for each of the three selected models.

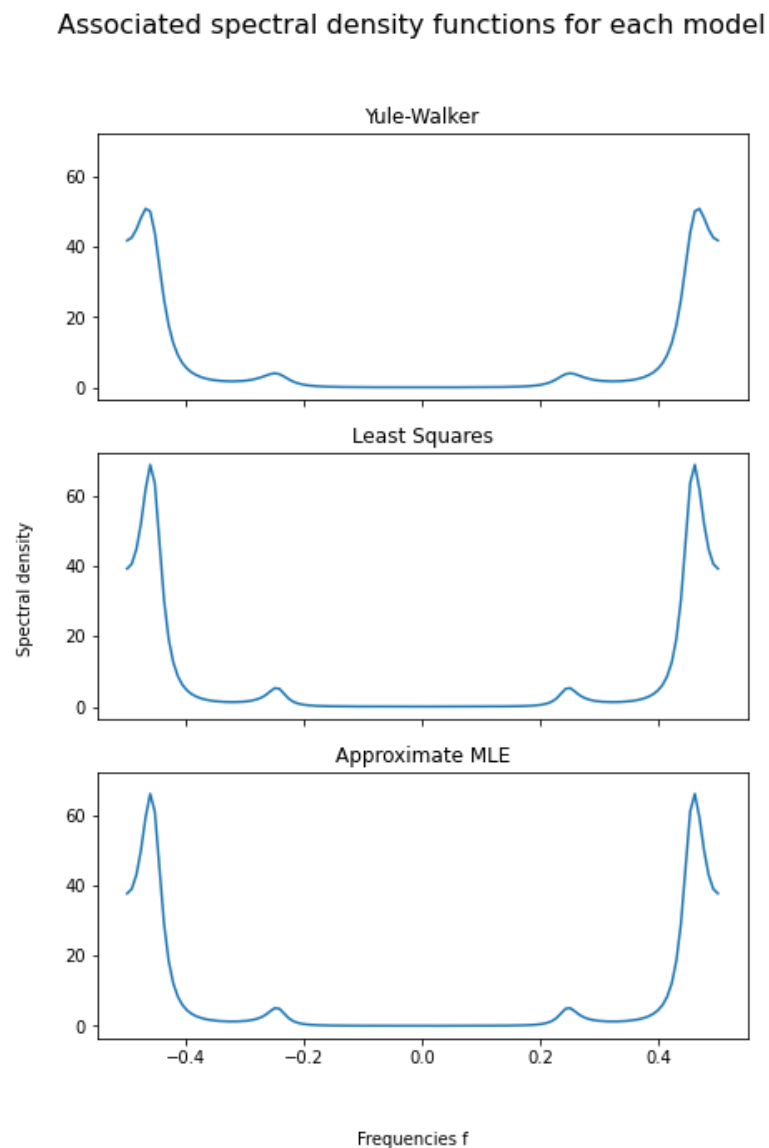


Figure 4: Plot of spectral densities for question 3(e)

```

1 def question 3e(X):
2     """
3     Code for question 3e
4     INPUT:
5         X: my time series
6     OUTPUT:
7         py, pl, pm: phi parameters for each method
8     """
9     #Compute fitted parameters for chosen p
10    py, sy = yule_walker(X,4)
11    pl,sl = least_squares(X, 5)
12    pm,sm = approx_mle(X, 5)
13    #Compute spectral densities
14    f = [i/128 for i in range(-64,65)]
15    S_y = S_AR(f, py, sy)
16    S_l = S_AR(f, pl, sl)
17    S_m = S_AR(f, pm, sm)
18    #plots
19    fig, ax = plt.subplots(3, 1, sharex = True, sharey = True, figsize = (7,10))
20    ax[0].plot(f, S_y, label = 'Yule-Walker')
21    ax[0].set_title('Yule-Walker')
22    ax[1].plot(f, S_l, label = 'Least Squares')
23    ax[1].set_title('Least Squares')
24    ax[2].plot(f, S_m, label = 'Approximate MLE')
25    ax[2].set_title('Approximate MLE')
26    fig.suptitle('Associated spectral density functions for each model',
27                size = 16)
28    fig.text(0.5, 0.04, 'Frequencies f', va='center', ha='center')
29    fig.text(0.04, 0.5, 'Spectral density', va='center', ha='center',
30             rotation='vertical')
31    plt.show()
32    return py, pl, pm

```

Question 4

The table below compares the actual values X_{19}, \dots, X_{128} to the forecast values, using the model parameters obtained from YW, LS and ML in question 3. I used the formula

$$X_t(l) = \phi_{1,p}X_t + \dots\phi_{p,p}X_{t-p+1}$$

from lecture notes to calculate each AR(p) forecast.

	Actual values	Yule-Walker	Least Squares	Approximate ML
t = 110	-3.73160	-3.731600	-3.731600	-3.731600
t = 111	2.14740	2.147400	2.147400	2.147400
t = 112	-1.69030	-1.690300	-1.690300	-1.690300
t = 113	2.23870	2.238700	2.238700	2.238700
t = 114	-2.21710	-2.217100	-2.217100	-2.217100
t = 115	1.72930	1.729300	1.729300	1.729300
t = 116	-1.18780	-1.187800	-1.187800	-1.187800
t = 117	0.26944	0.269440	0.269440	0.269440
t = 118	1.76370	1.763700	1.763700	1.763700
t = 119	-3.49900	-2.630999	-2.913407	-2.913407
t = 120	3.66870	2.065189	2.548070	2.548070
t = 121	-4.21590	-1.814666	-2.315649	-2.315649
t = 122	4.40100	2.111316	2.707106	2.707106
t = 123	-3.80270	-1.864365	-2.521683	-2.521683
t = 124	2.50420	1.102283	1.580764	1.580764
t = 125	0.36910	-0.685727	-0.964388	-0.964388
t = 126	-2.37090	0.619468	0.856305	0.856305
t = 127	2.95990	-0.354670	-0.487489	-0.487489
t = 128	-3.26850	-0.048823	-0.171230	-0.171230

Plot of time series X at time points 110 to 128, with actual values vs forecasted values from YS, LS and ML methods




Figure 5: Plot of forecast values for question 4

Figure 5 shows the plot of the forecast values against the true values for time points 110 to 128. Because the Least Squares and approximate ML methods give the same estimates of $\phi_{i,p}$, the forecast values are the same so the plots overlap.

```

1 def question4(X):
2     """
3     Code for question 4
4     INPUT:
5         X: my time series
6     OUTPUT:
7         data_f: dataframe of actual and forecast values at time points 110 to 128
8     """
9     #Get fitted parameters
10    py, pl, pm = question3e(X)
11    #Create matrix to contain actual and forecast values
12    X_f = np.zeros((128,4))
13    X_f[:,0] = X
14    X_f[:,118,1] = X[:,118]
15    X_f[:,118,2] = X[:,118]
16    X_f[:,118,3] = X[:,118]
17    for i in range(9):
18        #Loop to calculate forecast values for each method
19        X_f[118 + i,1] = py[0]*X_f[118 + i-1,1] +py[1]*X_f[118 + i-2,1] \
20            + py[2]*X_f[118 + i-3,1]+ py[3]*X_f[118 + i-4,1]
21        X_f[118 + i,2] = pl[0]*X_f[118 + i-1,2] +pl[1]*X_f[118 + i-2,2] \
22            + pl[2]*X_f[118 + i-3,2]+ pl[3]*X_f[118 + i-4,2] \
23            + pl[4]*X_f[118 + i-5,2]
24        X_f[118 + i,3] = pm[0]*X_f[118 + i-1,3] +pm[1]*X_f[118 + i-2,3] \
25            + pm[2]*X_f[118 + i-3,3]+ pm[3]*X_f[118 + i-4,3] \
26            + pm[4]*X_f[118 + i-5,3]
27    #take values only from time point 110
28    X_f = X_f[109,:,:]
29    #Plots
30    t = [i for i in range(110, 129)]
31    plt.figure(figsize = (10,7))
32    plt.plot(t, X_f[:,0], label = 'True values', marker = 'o')
33    plt.plot(t, X_f[:,1], label = 'Yule-Walker forecast', marker = 'o')
34    plt.plot(t, X_f[:,2], label = 'Least Squares forecast', marker = 'o')
35    plt.plot(t, X_f[:,3], label = 'Approximate ML forecast', marker = 'o')
36    plt.legend()

```

```
37 plt.title('Plot of time series X at time points 110 to 128, with actual '+
38           'values vs forecasted values from YS, LS and ML methods')
39 plt.xlabel('Time point t')
40 plt.ylabel('X_t / X_t(1)')
41 plt.show()
42 #summarise forecast and actual values into pandas Data Frame
43 data_f = pd.DataFrame(X_f, index = ['t = '+str(i) for i in range(110,129)],
44                        columns = ['Actual values', 'Yule-Walker',
45                                'Least Squares', ' Approximate ML'])
46 return data_f
```
